



# CMSC 5743

## Efficient Computing of Deep Neural Networks

### Mo03: Quantization

Bei Yu

CSE Department, CUHK

[byu@cse.cuhk.edu.hk](mailto:byu@cse.cuhk.edu.hk)

(Latest update: September 2, 2024)

2024 Fall



These slides contain/adapt materials developed by

- Hardware for Machine Learning, Shao Spring 2020 @ UCB
- 8-bit Inference with TensorRT
- Amir Gholami et al. (2021). “A survey of quantization methods for efficient neural network inference”. In: *arXiv preprint*



- ① Floating Point Number
- ② Integer & Fixed-Point Number
- ③ Quantization Overview
- ④ Quantization – First Example
- ⑤ Post Training Quantization (PTQ)
- ⑥ Quantization Aware Training (QAT)



# Floating Point Number



Scientific notation:  $6.6254 \times 10^{-27}$

- A normalized number of certain accuracy (e.g. 6.6254 is called the **mantissa**)
- Scale factors to determine the position of the decimal point (e.g.  $10^{-27}$  indicates position of decimal point and is called the exponent; the **base** is implied)
- **Sign bit**



- Floating Point Numbers can have multiple forms, e.g.


$$\begin{aligned}0.232 \times 10^4 &= 2.32 \times 10^3 \\&= 23.2 \times 10^2 \\&= 2320. \times 10^0 \\&= 232000. \times 10^{-2}\end{aligned}$$

- It is desirable for each number to have a unique representation => **Normalized Form**
- We normalize Mantissa's in the Range  $[1..R)$ , where R is the Base, e.g.:
  - [1..2) for BINARY
  - [1..10) for DECIMAL


# IEEE Standard 754 Single Precision



32-bit, float in C / C++ / Java



(a) Single precision



(b) Example of a single-precision number



64-bit, float in C / C++ / Java



(c) Double precision



## Question:

What is the IEEE single precision number  $40C0\ 0000_{16}$  in decimal?



## Question:

What is  $-0.5_{10}$  in IEEE single precision binary floating point format?



Exponents of all 0's and all 1's have special meaning

- E=0, M=0 represents 0 (sign bit still used so there is  $\pm 0$ )
- E=0, M $\neq$ 0 is a denormalized number  $\pm 0.M \times 2^{-126}$  (smaller than the smallest normalized number)
- E>All 1's, M=0 represents  $\pm \text{Infinity}$ , depending on Sign
- E>All 1's, M $\neq$ 0 represents NaN




- Normalized  $+/- 1.d\dots d \times 2^{\text{exp}}$
- Denormalized**  $+/- 0.d\dots d \times 2^{\min_{\text{exp}}}$  → to represent near-zero numbers  
e.g.  $+ 0.0000\dots 0000001 \times 2^{-126}$  for Single Precision

| Format          | # bits    | # significant bits | macheps                        | # exponent bits | exponent range                                 |
|-----------------|-----------|--------------------|--------------------------------|-----------------|--|
| Single          | 32        | 1+23               | $2^{-24} (\sim 10^{-7})$       | 8               | $2^{-126} - 2^{+127} (\sim 10^{\pm 38})$       |
| Double          | 64        | 1+52               | $2^{-53} (\sim 10^{-16})$      | 11              | $2^{-1022} - 2^{+1023} (\sim 10^{\pm 308})$    |
| Double Extended | $\geq 80$ | $\geq 64$          | $\leq 2^{-64} (\sim 10^{-19})$ | $\geq 15$       | $2^{-16382} - 2^{+16383} (\sim 10^{\pm 4932})$ |

(Double Extended is 80 bits on all Intel machines)

macheps = Machine Epsilon =  $2^{-(\# \text{ significant bits})}$



# Inaccurate Floating Point Operations



Example: Find 1st root of a quadratic equation<sup>1</sup>

$$r = \frac{-b + \sqrt{b^2 - 4 \cdot a \cdot c}}{2 \cdot a}$$

**Expected:** 0.00023025562642476431

**Double:** 0.00023025562638524986

**Float:** 0.00024670246057212353

<sup>1</sup>On Sparc processor, Solaris, gcc 3.3 (ANSI C)



Example: Find 1st root of a quadratic equation<sup>1</sup>

$$r = \frac{-b + \sqrt{b^2 - 4 \cdot a \cdot c}}{2 \cdot a}$$

**Expected:** 0.00023025562642476431

**Double:** 0.00023025562638524986

**Float:** 0.00024670246057212353

- Problem is that if  $c$  is near zero,  $\sqrt{b^2 - 4 \cdot a \cdot c} \approx b$
- **Rule of thumb:** use the highest precision which does not give up too much speed

<sup>1</sup>On Sparc processor, Solaris, gcc 3.3 (ANSI C)



# Integer & Fixed-Point Number

# Unsigned Binary Representation



| Hex        | Binary   | Decimal      |
|------------|----------|--------------|
| 0x00000000 | 0...0000 | 0            |
| 0x00000001 | 0...0001 | 1            |
| 0x00000002 | 0...0010 | 2            |
| 0x00000003 | 0...0011 | 3            |
| 0x00000004 | 0...0100 | 4            |
| 0x00000005 | 0...0101 | 5            |
| 0x00000006 | 0...0110 | 6            |
| 0x00000007 | 0...0111 | 7            |
| 0x00000008 | 0...1000 | 8            |
| 0x00000009 | 0...1001 | 9            |
|            | ...      |              |
| 0xFFFFFFF0 | 1...1100 | $2^{32} - 4$ |
| 0xFFFFFFF1 | 1...1101 | $2^{32} - 3$ |
| 0xFFFFFFF2 | 1...1110 | $2^{32} - 2$ |
| 0xFFFFFFFF | 1...1111 | $2^{32} - 1$ |

$2^{31}$   $2^{30}$   $2^{29}$  ...  $2^3$   $2^2$   $2^1$   $2^0$  bit weight

31 30 29 ... 3 2 1 0 bit position

1 1 1 ... 1 1 1 1 bit



1 0 0 0 ... 0 0 0 0 - 1



$2^{32} - 1$

# Signed Binary Representation



| 2'sc binary | decimal |
|-------------|---------|
| 1000        | -8      |
| 1001        | -7      |
| 1010        | -6      |
| 1011        | -5      |
| 1100        | -4      |
| 1101        | -3      |
| 1110        | -2      |
| 1111        | -1      |
| 0000        | 0       |
| 0001        | 1       |
| 0010        | 2       |
| 0011        | 3       |
| 0100        | 4       |
| 0101        | 5       |
| 0110        | 6       |
| 0111        | 7       |

$-2^3 =$

$-(2^3 - 1) =$

complement all the bits

0101      1011

and add a 1

0110      1010

complement all the bits

$2^3 - 1 =$

# Fixed-Point Arithmetic



- Integers with a binary point and a bias
  - “slope and bias”:  $y = s^*x + z$
  - $Qm.n$ : m (# of integer bits) n (# of fractional bits)

$$s = 1, z = 0$$

| 2^2 | 2^1 | 2^0 | Val |
|-----|-----|-----|-----|
| 0   | 0   | 0   | 0   |
| 0   | 0   | 1   | 1   |
| 0   | 1   | 0   | 2   |
| 0   | 1   | 1   | 3   |
| 1   | 0   | 0   | 4   |
| 1   | 0   | 1   | 5   |
| 1   | 1   | 0   | 6   |
| 1   | 1   | 1   | 7   |

$$s = 1/4, z = 0$$

| 2^0 | 2^-1 | 2^-2 | Val |
|-----|------|------|-----|
| 0   | 0    | 0    | 0   |
| 0   | 0    | 1    | 1/4 |
| 0   | 1    | 0    | 2/4 |
| 0   | 1    | 1    | 3/4 |
| 1   | 0    | 0    | 1   |
| 1   | 0    | 1    | 5/4 |
| 1   | 1    | 0    | 6/4 |
| 1   | 1    | 1    | 7/4 |

$$s = 4, z = 0$$

| 2^4 | 2^3 | 2^2 | Val |
|-----|-----|-----|-----|
| 0   | 0   | 0   | 0   |
| 0   | 0   | 1   | 4   |
| 0   | 1   | 0   | 8   |
| 0   | 1   | 1   | 12  |
| 1   | 0   | 0   | 16  |
| 1   | 0   | 1   | 20  |
| 1   | 1   | 0   | 24  |
| 1   | 1   | 1   | 28  |

$$s = 1.5, z = 10$$

| 2^2 | 2^1 | 2^0 | Val          |
|-----|-----|-----|--------------|
| 0   | 0   | 0   | $1.5*0 + 10$ |
| 0   | 0   | 1   | $1.5*1 + 10$ |
| 0   | 1   | 0   | $1.5*2 + 10$ |
| 0   | 1   | 1   | $1.5*3 + 10$ |
| 1   | 0   | 0   | $1.5*4 + 10$ |
| 1   | 0   | 1   | $1.5*5 + 10$ |
| 1   | 1   | 0   | $1.5*6 + 10$ |
| 1   | 1   | 1   | $1.5*7 + 10$ |



$(a - b)$  is inaccurate when  $a \gg b$  or  $a \ll b$

## Decimal Example 1:

- Using 2 significant digits
- Compute mean of 5.1 and 5.2 using the formula  $(a + b)/2$ :
- $a + b = 10$  (with 2 significant digits, 10.3 can only be stored as 10)
- $10/2 = 5.0$  (the computed mean is less than both numbers!!!)

## Decimal Example 2:

- Using 8 significant digits to compute sum of three numbers:
- $(11111113 + (-11111111)) + 7.5111111 = 9.5111111$
- $11111113 + ((-11111111) + 7.5111111) = 10.000000$



Catastrophic cancellation occurs when


$$\left| \frac{[\text{round}(x) \times \text{round}(y)] - \text{round}(x \times y)}{\text{round}(x \times y)} \right| >> \epsilon$$

# Hardware Implications




## Multipliers

Multiplier Example:  $C = A \times B$



Floating-point multiplier




Fixed-point multiplier



## Fixed-Point Arithmetic

*Number representation*  $\langle \text{IL}, \text{FL} \rangle$



Word Length  $WL = \text{IL} + \text{FL}$

Granularity  $2^{-\text{FL}}$

Range  $[-2^{\text{IL}-1}, 2^{\text{IL}-1} - 2^{-\text{FL}}]$

Convert  $(x, \langle \text{IL}, \text{FL} \rangle) =$

$$\begin{cases} -2^{\text{IL}-1} & \text{if } x \leq -2^{\text{IL}-1} \\ 2^{\text{IL}-1} - 2^{-\text{FL}} & \text{if } x \geq 2^{\text{IL}-1} - 2^{-\text{FL}} \\ \text{Round}(x, \langle \text{IL}, \text{FL} \rangle) & \text{otherwise} \end{cases}$$




<sup>2</sup>Suyog Gupta et al. (2015). “Deep learning with limited numerical precision”. In: *Proc. ICML*, pp. 1737–1746.



## Fixed-Point Arithmetic

*Number representation*  $\langle \text{IL}, \text{FL} \rangle$



$$\text{Word Length } \text{WL} = \text{IL} + \text{FL}$$

$$\text{Granularity } 2^{-\text{FL}}$$


$$\text{Range } [-2^{\text{IL}-1}, 2^{\text{IL}-1} - 2^{-\text{FL}}]$$

*Convert* ( $x, \langle \text{IL}, \text{FL} \rangle$ ) =

$$\begin{cases} -2^{\text{IL}-1} & \text{if } x \leq -2^{\text{IL}-1} \\ 2^{\text{IL}-1} - 2^{-\text{FL}} & \text{if } x \geq 2^{\text{IL}-1} - 2^{-\text{FL}} \\ \text{Round}(x, \langle \text{IL}, \text{FL} \rangle) & \text{otherwise} \end{cases}$$

8



*Multiply-and-ACCumulate*



<sup>2</sup>Suyog Gupta et al. (2015). “Deep learning with limited numerical precision”. In: *Proc. ICML*, pp. 1737–1746.





## Fixed-Point Arithmetic: Rounding Modes

*Round-to-nearest*




## Fixed-Point Arithmetic: Rounding Modes

*Round-to-nearest*



*Stochastic rounding*





$\text{Round}(x, \langle \text{IL}, \text{FL} \rangle) =$



$$\begin{cases} \lfloor x \rfloor & \text{w.p. } 1 - \frac{x - \lfloor x \rfloor}{\epsilon} \\ \lfloor x \rfloor + \epsilon & \text{w.p. } \frac{x - \lfloor x \rfloor}{\epsilon} \end{cases}$$

- Non-zero probability of rounding to either  $\lfloor x \rfloor$  or  $\lfloor x \rfloor + \epsilon$
- Unbiased rounding scheme: expected rounding error is zero



MNIST: *Fully-connected DNNs*



<sup>2</sup>Suyog Gupta et al. (2015). “Deep learning with limited numerical precision”. In: *Proc. ICML*, pp. 1737–1746.

MNIST: *Fully-connected DNNs*

- For small fractional lengths ( $FL < 12$ ), a large majority of weight updates are rounded to zero when using the round-to-nearest scheme.
  - Convergence slows down
- For  $FL < 12$ , there is a noticeable degradation in the classification accuracy



<sup>2</sup>Suyog Gupta et al. (2015). “Deep learning with limited numerical precision”. In: *Proc. ICML*, pp. 1737–1746.

MNIST: *Fully-connected DNNs*


- Stochastic rounding preserves gradient information (statistically)
  - No degradation in convergence properties
- Test error nearly equal to that obtained using 32-bit floats



<sup>2</sup>Suyog Gupta et al. (2015). “Deep learning with limited numerical precision”. In: *Proc. ICML*, pp. 1737–1746.




## FPGA prototyping: GEMM with stochastic rounding



<sup>2</sup>Suyog Gupta et al. (2015). “Deep learning with limited numerical precision”. In: *Proc. ICML*, pp. 1737–1746.





## Stochastic rounding





## Stochastic rounding



These operations can be implemented efficiently using a single DSP unit



<sup>2</sup>Suyog Gupta et al. (2015). "Deep learning with limited numerical precision". In: *Proc. ICML*, pp. 1737–1746.



# Quantization Overview

# Quantization in DNN



## Quantization:


$$Q(r) = \text{Int}(r/S) - Z$$

## Dequantization:


$$\hat{r} = S(Q(r) + Z)$$

## Granularity:

- Layerwise
- Groupwise
- Channelwise




# Uniform vs. Non-Uniform




- Real values in the continuous domain  $r$  are mapped into discrete
- Lower precision values in the quantized domain  $Q$ .
- **Uniform** quantization: distances between quantized values are **the same**
- **Non-uniform** quantization: distances between quantized values can **vary**

# Symmetric vs. Asymmetric




(a) Symmetric quantization




(b) Asymmetric quantization

- Symmetric vs. Asymmetric:  $Z = 0$  ?
- Fig. (a) Symmetric w. **restricted range** maps  $[-127, 127]$ ,
- Fig. (b) Asymmetric w. **full range** maps to  $[-128, 127]$
- Both for 8-bit quantization case.



- **quantization-aware training (QAT)**: model is quantized using training data to adjust parameters and recover accuracy degradation.
- **post-training quantization (PTQ)**: a pre-trained model is calibrated using finetuning data (e.g., a small subset of training data) to compute the clipping ranges and the scaling factors.
- **Key difference**: Model parameters fixed/unfixed.

# Simulated quantization vs Integer-Only quantization



**Left** : Full-precision

**Middle** : Simulated quantization

**Right** : Integer-only quantization



## Hardware Support

- Nvidia GPU: Tensor Core support FP16, Int8 and Int4
- Arm: Neon 128-bit SIMD instruction:  $4 \times 32\text{bit}$  or  $8 \times 16\text{bit}$  up to  $16 \times 8\text{bit}$
- Intel: SSE intrinsics, same as above
- DSP, AI Chip

## Some common architectures:

- For CPU: Tensorflow Lite, QNNPACK, NCNN
- For GPU: TensorRT
- Versatile Compiler such TVM.qnn



# Quantization – First Example



# Linear quantization

Representation:

Tensor Values = FP32 scale factor \* int8 array + FP32 bias



# Do we really need bias?

Two matrices:

```
A = scale_A * QA + bias_A  
B = scale_B * QB + bias_B
```

Let's multiply those 2 matrices:

```
A * B = scale_A * scale_B * QA * QB +  
        scale_A * QA * bias_B +  
        scale_B * QB * bias_A +  
        bias_A * bias_B
```



# Do we really need bias?

Two matrices:

```
A = scale_A * QA + bias_A  
B = scale_B * QB + bias_B
```

Let's multiply those 2 matrices:

```
A * B = scale_A * scale_B * QA * QB +  
        scale_A * QA * bias_B      +  
        scale_B * QB * bias_A      +  
        bias_A * bias_B
```



# Do we really need bias? No!

Two matrices:

```
A = scale_A * QA  
B = scale_B * QB
```

Let's multiply those 2 matrices:

```
A * B = scale_A * scale_B * QA * QB
```



# Symmetric linear quantization

Representation:

Tensor Values = **FP32 scale factor** \* int8 array

One **FP32 scale factor** for the entire int8 tensor

**Q: How do we set scale factor?**



# MINIMUM QUANTIZED VALUE

- Integer range is not completely symmetric. E.g. in 8bit, [-128, 127]
  - If use [-127, 127],  $s = \frac{127}{\alpha}$ 
    - Range is symmetric
    - 1/256 of int8 range is not used. 1/16 of int4 range is not used
  - If use full range [-128, 127],  $s = \frac{128}{\alpha}$ 
    - Values should be quantized to 128 will be clipped to 127
    - Asymmetric range may introduce bias



# EXAMPLE OF QUANTIZATION BIAS

Bias introduced when int values are in [-128, 127]

$$A = [-2.2 \quad -1.1 \quad 1.1 \quad 2.2], B = \begin{bmatrix} 0.5 \\ 0.3 \\ 0.3 \\ 0.5 \end{bmatrix}, AB = 0$$

8bit scale quantization, use [-128, 127].  $s_A=128/2.2$ ,  $s_B=128/0.5$

$$[-128 \quad -64 \quad 64 \quad 127] * \begin{bmatrix} 127 \\ 77 \\ 77 \\ 127 \end{bmatrix} = -127$$

Dequantize -127 will get -0.00853. A small bias is introduced towards  $-\infty$



# EXAMPLE OF QUANTIZATION BIAS

No bias when int values are in [-127, 127]

$$A = [-2.2 \quad -1.1 \quad 1.1 \quad 2.2], B = \begin{bmatrix} 0.5 \\ 0.3 \\ 0.3 \\ 0.5 \end{bmatrix}, AB = 0$$

8-bit scale quantization, use [-127, 127].  $s_A=127/2.2$ ,  $s_B=127/0.5$

$$[-127 \quad -64 \quad 64 \quad 127] * \begin{bmatrix} 127 \\ 76 \\ 76 \\ 127 \end{bmatrix} = 0$$

Dequantize 0 will get 0



# MATRIX MULTIPLY EXAMPLE

## Scale Quantization

$$\begin{pmatrix} -1.54 & 0.22 \\ -0.26 & 0.65 \end{pmatrix} * \begin{pmatrix} 0.35 \\ -0.51 \end{pmatrix} = \begin{pmatrix} -0.651 \\ -0.423 \end{pmatrix}$$



# MATRIX MULTIPLY EXAMPLE

## Scale Quantization

$$\begin{pmatrix} -1.54 & 0.22 \\ -0.26 & 0.65 \end{pmatrix} * \begin{pmatrix} 0.35 \\ -0.51 \end{pmatrix} = \begin{pmatrix} -0.651 \\ -0.423 \end{pmatrix}$$

8bit quantization

choose [-2, 2] fp range (scale 127/2=63.5) for first matrix and [-1, 1] fp range (scale = 127/1=127) for the second

$$\begin{pmatrix} -98 & 14 \\ -17 & 41 \end{pmatrix} * \begin{pmatrix} 44 \\ -65 \end{pmatrix} = \begin{pmatrix} -5222 \\ -3413 \end{pmatrix}$$



# MATRIX MULTIPLY EXAMPLE

## Scale Quantization

$$\begin{pmatrix} -1.54 & 0.22 \\ -0.26 & 0.65 \end{pmatrix} * \begin{pmatrix} 0.35 \\ -0.51 \end{pmatrix} = \begin{pmatrix} -0.651 \\ -0.423 \end{pmatrix}$$

8bit quantization

choose [-2, 2] fp range (scale  $127/2=63.5$ ) for first matrix and [-1, 1] fp range (scale =  $127/1=127$ ) for the second

$$\begin{pmatrix} -98 & 14 \\ -17 & 41 \end{pmatrix} * \begin{pmatrix} 44 \\ -65 \end{pmatrix} = \begin{pmatrix} -5222 \\ -3413 \end{pmatrix}$$

The result has an overall scale of  $63.5 * 127$ . We can dequantize back to float

$$\begin{pmatrix} -5222 \\ -3413 \end{pmatrix} * \frac{1}{63.5 * 127} = \begin{pmatrix} -0.648 \\ -0.423 \end{pmatrix}$$



# REQUANTIZE

## Scale Quantization

$$\begin{pmatrix} -1.54 & 0.22 \\ -0.26 & 0.65 \end{pmatrix} * \begin{pmatrix} 0.35 \\ -0.51 \end{pmatrix} = \begin{pmatrix} -0.651 \\ -0.423 \end{pmatrix}$$

8bit quantization

choose [-2, 2] fp range for first matrix and [-1, 1] fp range for the second

$$\begin{pmatrix} -98 & 14 \\ -17 & 41 \end{pmatrix} * \begin{pmatrix} 44 \\ -65 \end{pmatrix} = \begin{pmatrix} -5222 \\ -3413 \end{pmatrix}$$

*Requantize* output to a different quantized representation with fp range [-3, 3]:

$$\begin{pmatrix} -5222 \\ -3413 \end{pmatrix} * \frac{127/3}{63.5 * 127} = \begin{pmatrix} -27 \\ -18 \end{pmatrix}$$



# Post Training Quantization (PTQ)



- For a fixed-point number, its representation is:

$$n = \sum_{i=0}^{bw-1} B_i \cdot 2^{-f_l} \cdot 2^i,$$

where  $bw$  is the bit width and  $f_l$  is the fractional length which is dynamic for different layers and feature map sets while static in one layer.

- Weight quantization: find the optimal  $f_l$  for weights:

$$f_l = \arg \min_{f_l} \sum |W_{\text{float}} - W(bw, f_l)|,$$

where  $W$  is a weight and  $W(bw, f_l)$  represents the fixed-point format of  $W$  under the given  $bw$  and  $f_l$ .

---

<sup>3</sup>Jiantao Qiu et al. (2016). "Going deeper with embedded fpga platform for convolutional neural network". In: *Proc. FPGA*, pp. 26–35.


# Greedy Layer-wise Quantization



- Feature quantization: find the optimal  $f_l$  for features:

$$f_l = \arg \min_{f_l} \sum |x_{\text{float}}^+ - x^+(bw, f_l)|,$$

where  $x^+$  represents the result of a layer when we denote the computation of a layer as  $x^+ = A \cdot x$ .




# Dynamic-Precision Data Quantization Results



| Network          | VGG16        |           |          |              |                 |         |         |
|------------------|--------------|-----------|----------|--------------|-----------------|---------|---------|
| Data Bits        | Single-float | 16        | 16       | 8            | 8               | 8       | 8       |
| Weight Bits      | Single-float | 16        | 8        | 8            | 8               | 8       | 8 or 4  |
| Data Precision   | N/A          | $2^{-2}$  | $2^{-2}$ | Impossible   | $2^{-5}/2^{-1}$ | Dynamic | Dynamic |
| Weight Precision | N/A          | $2^{-15}$ | $2^{-7}$ | Impossible   | $2^{-7}$        | Dynamic | Dynamic |
| Top-1 Accuracy   | 68.1%        | 68.0%     | 53.0%    | Impossible   | 28.2%           | 66.6%   | 67.0%   |
| Top-5 Accuracy   | 88.0%        | 87.9%     | 76.6%    | Impossible   | 49.7%           | 87.4%   | 87.6%   |
| Network          | CaffeNet     |           |          | VGG16-SVD    |                 |         |         |
| Data Bits        | Single-float | 16        | 8        | Single-float | 16              | 8       |         |
| Weight Bits      | Single-float | 16        | 8        | Single-float | 16              | 8 or 4  |         |
| Data Precision   | N/A          | Dynamic   | Dynamic  | N/A          | Dynamic         | Dynamic |         |
| Weight Precision | N/A          | Dynamic   | Dynamic  | N/A          | Dynamic         | Dynamic |         |
| Top-1 Accuracy   | 53.9%        | 53.9%     | 53.0%    | 68.0%        | 64.6%           | 64.1%   |         |
| Top-5 Accuracy   | 77.7%        | 77.1%     | 76.6%    | 88.0%        | 86.7%           | 86.3%   |         |




## No Saturation Quantization – INT8 Inference



- Map the maximum value to 127, with uniform step length.
- Suffer from outliers.



## Saturation Quantization – INT8 Inference



- Set a threshold as the maximum value.
- Divide the value domain into 2048 groups.
- Traverse all the possible thresholds to find the best one with minimum KL divergence.



## Relative Entropy of two encodings

- INT8 model encodes the same information as the original FP32 model.
- Minimize the loss of information.
- Loss of information is measured by **Kullback-Leibler divergence** (*a.k.a.*, relative entropy or information divergence).
  - $P, Q$  - two discrete probability distributions:

$$D_{KL}(P\|Q) = \sum_{i=1}^N P(x_i) \log \frac{P(x_i)}{Q(x_i)}$$

- Intuition: KL divergence measures **the amount of information lost** when approximating a given encoding.




# Quantization Aware Training (QAT)



## Straight Through Estimator (STE)<sup>4</sup>

- Forward integer, Backward floating point
- Rounding to nearest



<sup>4</sup>Yoshua Bengio, Nicholas Léonard, and Aaron Courville (2013). "Estimating or propagating gradients through stochastic neurons for conditional computation". In: *arXiv preprint arXiv:1308.3432*.



## Is Straight-Through Estimator (STE) the best?


- Gradient mismatch: the gradients of the weights are not generated using the value of weights, but rather its quantized value.
- Poor gradient: STE fails at investigating better gradients for quantization training.



## PParameterized Clipping acTivation (PACT)<sup>5</sup>

- Relu6 → clipping
- threshold → clipping range in quantization
- range upper/lower bound trainable

$$y = PACT(x) = 0.5(|x| - |x - \alpha| + \alpha) = \begin{cases} 0, & x \in (-\infty, 0) \\ x, & x \in [0, \alpha) \\ \alpha, & x \in [\alpha, +\infty) \end{cases}$$



<sup>5</sup>Jungwook Choi, Zhuo Wang, et al. (2018). "Pact: Parameterized clipping activation for quantized neural networks". In: *arXiv preprint arXiv:1805.06085*.



- A new activation quantization scheme in which the activation function has a parameterized clipping level  $\alpha$ .
- The clipping level is dynamically adjusted via stochastic gradient descent (SGD)-based training with the goal of minimizing the quantization error.
- In PACT, the convolutional ReLU activation function in CNN is replaced with:

$$f(x) = 0.5 (|x| - |x - \alpha| + \alpha) = \begin{cases} 0, & x \in (\infty, 0) \\ x, & x \in [0, \alpha) \\ \alpha, & x \in [\alpha, +\infty) \end{cases}$$

where  $\alpha$  limits the dynamic range of activation to  $[0, \alpha]$ .

---

<sup>6</sup>Jungwook Choi, Swagath Venkataramani, et al. (2019). "Accurate and efficient 2-bit quantized neural networks". In: *Proceedings of Machine Learning and Systems* 1.


# Parameterized Clipping Activation Function (PACT)



- The truncated activation output is linearly quantized to  $k$ -bits for the dot-product computations:

$$y_q = \text{round} \left( y \cdot \frac{2^k - 1}{\alpha} \right) \cdot \frac{\alpha}{2^k - 1}$$

- With this new activation function,  $\alpha$  is a variable in the loss function, whose value can be optimized during training.
- For back-propagation, gradient  $\frac{\partial y_q}{\partial \alpha}$  can be computed using STE to estimate  $\frac{\partial y_q}{\partial y}$  as 1.



PACT activation function and its gradient.